

# Steven Maharaj 695281 Assignment 2, Question 3

## MAST90125: Bayesian Statistical Learning

**Due: Friday 20 September 2019**

There are places in this assignment where R code will be required. Therefore set the random seed so assignment is reproducible.

```
set.seed(695281) #Please change random seed to your student id number.
library(dplyr)

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(ggplot2)
library(tidyr)
library(TruncatedNormal)
```

### Question Three (18 marks)

A group of 453 Bangladeshi women in 5 districts were asked about contraceptive use. The response variable *use* is an indicator for contraceptive use (coded N for no and Y for yes). Other covariates of interest are categorical variables for geographical location *district* (5 levels), and *urban* (2 levels), and number of living children *livch* (4 levels), and the continuous covariate for standardised age *age*. A random intercept for the district was suggested. This suggested the following model should be fitted,

$$\theta = \mathbf{Z}\mathbf{u} + \mathbf{X}\boldsymbol{\beta},$$

where  $\theta$  is a link function,  $\mathbf{Z}$  is an indicator variable for district,  $\mathbf{u}$  is a random intercept with prior  $p(\mathbf{u}) = \mathcal{N}(\mathbf{0}, \sigma_u^2 \mathbf{I})$ , and  $\mathbf{X}$  is a design matrix for fixed effects  $\boldsymbol{\beta}$ , where  $\boldsymbol{\beta}$  includes the coefficients for the intercept, urban status, living children, and age.

Data can be downloaded from LMS as `Contraceptionsubset.csv`.

- a) Fit a generalised linear mixed model assuming a logistic link using **Stan**. The R and stan code below covers the following steps.
  - Importing the data.
  - Constructing design matrices.
  - Provides code to go into the stan file.
  - Running stan in R. This assumes your stan file is called `*logitmm.stan*`, and that you will run the sampler for 2000 iterations and 4 chains.

Note that provided code assumes everything required is located in your working directory in R.

*#Step one: Importing data, constructing design matrices and calculating matrix dimensions.*  
dataX= read.csv("Contraceptionsubset.csv",header=TRUE)

```
n<-dim(dataX)[1]
Z  = table(1:n,dataX$district)      #incidence matrix for district
Q  = dim(Z)[2]
D1 = table(1:n,dataX$livch) #Dummy indicator for living children
D2 = table(1:n,dataX$urban) #Dummy indicator for urban status

#fixed effect design matrix
X  = cbind(rep(1,n),dataX$age,D1[,-1],D2[,-1])
P  = dim(X)[2]
y  = rep(0,n)
y[dataX$use %in% 'Y'] = 1
```

An example stan file.

```
//
// This Stan program defines a logistic mixed model
//
// Learn more about model development with Stan at:
//
//   http://mc-stan.org/users/interfaces/rstan.html
//   https://github.com/stan-dev/rstan/wiki/RStan-Getting-Started
//

data {
  int<lower=0> n;    //number of observations
  int<lower=0> Q;    //number of random effect levels
  int<lower=0> P;    //number of fixed effect levels
  int y[n];         //response vector
  matrix[n,Q] Z;    //indicator matrix for random effect levels
  matrix[n,P] X;    //design matrix for fixed effects
}

// The parameters accepted by the model.
// accepts three sets of parameters 'beta', 'u' and 'sigma'.
parameters {
  vector[P] beta; //vector of fixed effects of length P.
  vector[Q] u;    //vector of random effects of length Q.
  real<lower=0> sigma; //random effect standard deviation
}

// The model to be estimated. We model the output
// 'y' to be bernoulli with logit link function,
// and assume a i.i.d. normal prior for u.
model {
  u ~ normal(0,sigma);           //prior for random effects.
  y ~ bernoulli_logit(X*beta+ Z*u); //likelihood
}
```

```
library(rstan)
```

```
## Loading required package: StanHeaders
```

```
## rstan (Version 2.19.2, GitRev: 2e1f913d3ca3)
```

```
## For execution on a local, multicore CPU with excess RAM we recommend calling
## options(mc.cores = parallel::detectCores()).
```

```
## To avoid recompilation of unchanged Stan programs, we recommend calling
```

```

## rstan_options(auto_write = TRUE)

##
## Attaching package: 'rstan'

## The following object is masked from 'package:tidyr':
##
##      extract

logistic.mm <-stan(file="logitmm.stan",data=c('Z','X','y','n','P','Q'),iter=2000,chains=4)

##
## SAMPLING FOR MODEL 'logitmm' NOW (CHAIN 1).
## Chain 1:
## Chain 1: Gradient evaluation took 0.000104 seconds
## Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 1.04 seconds.
## Chain 1: Adjust your expectations accordingly!
## Chain 1:
## Chain 1:
## Chain 1: Iteration:    1 / 2000 [  0%] (Warmup)
## Chain 1: Iteration:   200 / 2000 [ 10%] (Warmup)
## Chain 1: Iteration:   400 / 2000 [ 20%] (Warmup)
## Chain 1: Iteration:   600 / 2000 [ 30%] (Warmup)
## Chain 1: Iteration:   800 / 2000 [ 40%] (Warmup)
## Chain 1: Iteration:  1000 / 2000 [ 50%] (Warmup)
## Chain 1: Iteration:  1001 / 2000 [ 50%] (Sampling)
## Chain 1: Iteration:  1200 / 2000 [ 60%] (Sampling)
## Chain 1: Iteration:  1400 / 2000 [ 70%] (Sampling)
## Chain 1: Iteration:  1600 / 2000 [ 80%] (Sampling)
## Chain 1: Iteration:  1800 / 2000 [ 90%] (Sampling)
## Chain 1: Iteration:  2000 / 2000 [100%] (Sampling)
## Chain 1:
## Chain 1: Elapsed Time: 1.53192 seconds (Warm-up)
## Chain 1:                    1.20015 seconds (Sampling)
## Chain 1:                    2.73207 seconds (Total)
## Chain 1:
##
## SAMPLING FOR MODEL 'logitmm' NOW (CHAIN 2).
## Chain 2:
## Chain 2: Gradient evaluation took 4.3e-05 seconds
## Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.43 seconds.
## Chain 2: Adjust your expectations accordingly!
## Chain 2:
## Chain 2:
## Chain 2: Iteration:    1 / 2000 [  0%] (Warmup)
## Chain 2: Iteration:   200 / 2000 [ 10%] (Warmup)
## Chain 2: Iteration:   400 / 2000 [ 20%] (Warmup)
## Chain 2: Iteration:   600 / 2000 [ 30%] (Warmup)
## Chain 2: Iteration:   800 / 2000 [ 40%] (Warmup)
## Chain 2: Iteration:  1000 / 2000 [ 50%] (Warmup)
## Chain 2: Iteration:  1001 / 2000 [ 50%] (Sampling)
## Chain 2: Iteration:  1200 / 2000 [ 60%] (Sampling)
## Chain 2: Iteration:  1400 / 2000 [ 70%] (Sampling)
## Chain 2: Iteration:  1600 / 2000 [ 80%] (Sampling)
## Chain 2: Iteration:  1800 / 2000 [ 90%] (Sampling)
## Chain 2: Iteration:  2000 / 2000 [100%] (Sampling)

```

```

## Chain 2:
## Chain 2: Elapsed Time: 1.61333 seconds (Warm-up)
## Chain 2: 1.25627 seconds (Sampling)
## Chain 2: 2.8696 seconds (Total)
## Chain 2:
##
## SAMPLING FOR MODEL 'logitmm' NOW (CHAIN 3).
## Chain 3:
## Chain 3: Gradient evaluation took 4.9e-05 seconds
## Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0.49 seconds.
## Chain 3: Adjust your expectations accordingly!
## Chain 3:
## Chain 3:
## Chain 3: Iteration: 1 / 2000 [ 0%] (Warmup)
## Chain 3: Iteration: 200 / 2000 [ 10%] (Warmup)
## Chain 3: Iteration: 400 / 2000 [ 20%] (Warmup)
## Chain 3: Iteration: 600 / 2000 [ 30%] (Warmup)
## Chain 3: Iteration: 800 / 2000 [ 40%] (Warmup)
## Chain 3: Iteration: 1000 / 2000 [ 50%] (Warmup)
## Chain 3: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 3: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 3: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 3: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 3: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 3: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 3:
## Chain 3: Elapsed Time: 1.50747 seconds (Warm-up)
## Chain 3: 1.17569 seconds (Sampling)
## Chain 3: 2.68315 seconds (Total)
## Chain 3:
##
## SAMPLING FOR MODEL 'logitmm' NOW (CHAIN 4).
## Chain 4:
## Chain 4: Gradient evaluation took 4.3e-05 seconds
## Chain 4: 1000 transitions using 10 leapfrog steps per transition would take 0.43 seconds.
## Chain 4: Adjust your expectations accordingly!
## Chain 4:
## Chain 4:
## Chain 4: Iteration: 1 / 2000 [ 0%] (Warmup)
## Chain 4: Iteration: 200 / 2000 [ 10%] (Warmup)
## Chain 4: Iteration: 400 / 2000 [ 20%] (Warmup)
## Chain 4: Iteration: 600 / 2000 [ 30%] (Warmup)
## Chain 4: Iteration: 800 / 2000 [ 40%] (Warmup)
## Chain 4: Iteration: 1000 / 2000 [ 50%] (Warmup)
## Chain 4: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 4: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 4: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 4: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 4: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 4: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 4:
## Chain 4: Elapsed Time: 2.12798 seconds (Warm-up)
## Chain 4: 1.21135 seconds (Sampling)
## Chain 4: 3.33933 seconds (Total)

```

```
## Chain 4:
```

```
print(logistic.mm)
```

```
## Inference for Stan model: logitmm.
```

```
## 4 chains, each with iter=2000; warmup=1000; thin=1;
```

```
## post-warmup draws per chain=1000, total post-warmup draws=4000.
```

```
##
```

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff
## beta[1]	-2.01	0.02	0.59	-3.21	-2.36	-2.00	-1.66	-0.81	1126
## beta[2]	-0.04	0.00	0.02	-0.08	-0.05	-0.04	-0.03	-0.01	2301
## beta[3]	1.23	0.01	0.34	0.57	1.01	1.23	1.46	1.89	2432
## beta[4]	1.45	0.01	0.36	0.74	1.20	1.44	1.69	2.15	2234
## beta[5]	1.79	0.01	0.38	1.07	1.54	1.78	2.04	2.54	1610
## beta[6]	1.22	0.00	0.26	0.72	1.04	1.22	1.39	1.74	2732
## u[1]	-1.06	0.02	0.56	-2.27	-1.37	-1.04	-0.72	-0.06	1082
## u[2]	-0.25	0.02	0.57	-1.41	-0.57	-0.26	0.08	0.83	1239
## u[3]	0.43	0.02	0.56	-0.76	0.12	0.43	0.74	1.50	1055
## u[4]	0.18	0.02	0.56	-0.95	-0.13	0.18	0.51	1.28	1078
## u[5]	0.69	0.02	0.56	-0.44	0.37	0.67	1.00	1.82	1197
## sigma	1.10	0.02	0.70	0.41	0.70	0.93	1.27	2.86	1480
## lp__	-274.34	0.07	2.51	-280.05	-275.89	-274.04	-272.49	-270.41	1246

```
## Rhat
```

## beta[1]	1
## beta[2]	1
## beta[3]	1
## beta[4]	1
## beta[5]	1
## beta[6]	1
## u[1]	1
## u[2]	1
## u[3]	1
## u[4]	1
## u[5]	1
## sigma	1
## lp__	1

```
## Samples were drawn using NUTS(diag_e) at Wed Sep 25 12:08:50 2019.
```

```
## For each parameter, n_eff is a crude measure of effective sample size,
```

```
## and Rhat is the potential scale reduction factor on split chains (at
```

```
## convergence, Rhat=1).
```

Note that in Stan, defaults for burn-in (warm-up) is one half of all iterations in stan, and no thinning. Note the code is written using the stan file and csv is in your working directory. Use the `print` function to report posterior means, standard deviations, 95 % central credible intervals and state from the output whether you believe the chains have converged. Also report the reference categories for *urban* and *livch*.

Reporting for PART A:

posterior means, standard deviations can be found in the above table. The lower limit of the 95 % central credible intervals given by the column labeled “2.5%” while the upper limit of the 95 % central credible intervals given by the column labeled “97.5%”.

```
print(summary(logistic.mm)$summary)
```

	mean	se_mean	sd	2.5%	25%
## beta[1]	-2.01030050	0.0177090209	0.59412704	-3.2126422	-2.35819525

```
## beta[2]    -0.04214017  0.0003647766  0.01749937   -0.0771671   -0.05338352
## beta[3]     1.23207517  0.0068231905  0.33646783    0.5659346    1.01390920
## beta[4]     1.44766009  0.0076461800  0.36137498    0.7375624    1.20459896
## beta[5]     1.78929396  0.0093728704  0.37613535    1.0743897    1.53649493
## beta[6]     1.21639026  0.0049852543  0.26059271    0.7185883    1.03977246
## u[1]        -1.06275576  0.0171555132  0.56436421   -2.2699659   -1.37142413
## u[2]        -0.25489581  0.0161806806  0.56952451   -1.4074344   -0.57137738
## u[3]         0.42541069  0.0173308629  0.56299322   -0.7568421    0.12006699
## u[4]         0.18475245  0.0169065649  0.55506518   -0.9452371   -0.12563814
## u[5]         0.69203489  0.0160977691  0.55701986   -0.4360102    0.36515485
## sigma       1.10279222  0.0181485968  0.69807840    0.4141299    0.69594181
## lp__        -274.33957141  0.0710899545  2.50946526 -280.0506863 -275.89450998
##              50%          75%          97.5%    n_eff    Rhat
## beta[1]     -2.00071542   -1.66048095  -8.137643e-01 1125.562  1.0028517
## beta[2]     -0.04190229   -0.03039232  -9.684808e-03 2301.393  1.0008164
## beta[3]      1.22785036    1.45582991   1.893473e+00 2431.712  1.0003286
## beta[4]      1.44472312    1.69041346   2.151146e+00 2233.712  1.0022645
## beta[5]      1.77969771    2.03909874   2.538047e+00 1610.435  1.0021622
## beta[6]      1.21564004    1.38717360   1.744403e+00 2732.435  1.0020165
## u[1]        -1.03749524   -0.72160839  -6.331785e-02 1082.210  1.0035870
## u[2]        -0.25724419    0.07818563   8.251233e-01 1238.886  1.0028778
## u[3]         0.43126008    0.74223850   1.496745e+00 1055.276  1.0032488
## u[4]         0.17671158    0.50561014   1.278456e+00 1077.897  1.0029502
## u[5]         0.66847636    1.00332512   1.815226e+00 1197.319  1.0031387
## sigma       0.93149252    1.27266697   2.858592e+00 1479.525  1.0007791
## lp__        -274.04470478 -272.48530390 -2.704118e+02 1246.080  0.9999782
```

Using the Gelman-Rubin diagnostic (Rhat) only 6 out of the 12 parameters These were beta[2],beta[3],beta[4],beta[5],beta[6],sigma had an Rhat value close to 1. All u parameters and beta[1] do not appear to converge thus the chain did not converge.

For *urban* and *livch* the reference categories are “N” and “0” respectively. This is inferred from the way the DataFrame X was constructed. `X = cbind(rep(1,n),dataX$age,D1[, -1],D2[, -1])`

b) An alternative to the logit link when analysing binary data is the probit. The probit link is defined as,

$$y_i = \begin{cases} 1 & \text{if } z_i \geq 0 \\ 0 & \text{if } z_i < 0 \end{cases}$$

$$z_i = \mathbf{x}_i' \boldsymbol{\beta} + \epsilon_i, \quad \epsilon \sim \mathcal{N}(0, 1).$$

In lecture 14, we showed how by letting  $z_i$  be normal, probit regression can be fitted using a Gibbs sampler, but to do so, it requires the ability to sample from a truncated normal defined on either  $(-\infty, 0)$  (if  $y_i = 0$ ) or  $(0, \infty)$  (if  $y_i = 1$ ). Check by comparing the empirical and the true density that a modified version of the inverse cdf method can be used to produce draws from a truncated normal. Do this for the case where  $x \in (0, \infty)$  and  $x \in (-\infty, 0)$  with parameters  $\mu = 0.5$  and  $\sigma = 1$ .

Hints: If  $y$  is drawn from a truncated normal with lower bound  $a$ , upper bound  $b$  and parameters  $\mu, \sigma^2$  then then  $p(y|\mu, \sigma^2, a, b)$  is

$$\frac{\frac{1}{\sqrt{2\pi\sigma^2}}e^{-(y-\mu)^2/2}}{\int_{-\infty}^b \frac{1}{\sqrt{2\pi\sigma^2}}e^{-(y-\mu)^2/2}dy - \int_{-\infty}^a \frac{1}{\sqrt{2\pi\sigma^2}}e^{-(y-\mu)^2/2}dy},$$

which in R means the truncated normal density can be written as

```
dnorm(x,mean=mu,sd=sigma)/(pnorm(b,mean=mu,sd=sigma)-pnorm(a,mean=mu,sd=sigma))
```

The inverse cdf method involves drawing  $v$  from  $U(0, 1)$  so that  $x \sim p(x)$  can be found solving  $x = F^{-1}(x)$ , where  $F$  is the cdf. If the only change compared to drawing from a normal distribution is truncation, think about what happens to the bounds of the uniform distribution.

Answer: Part B

Case  $x \in (0, \infty)$

```
rtn <- function(n,b,a,mu,Sigma){
  u <- runif(n)
  g <- pnorm((b-mu)/Sigma) - pnorm((a-mu)/Sigma)
  x <- qnorm((g) * u + pnorm((a-mu)/Sigma))*Sigma + mu
}

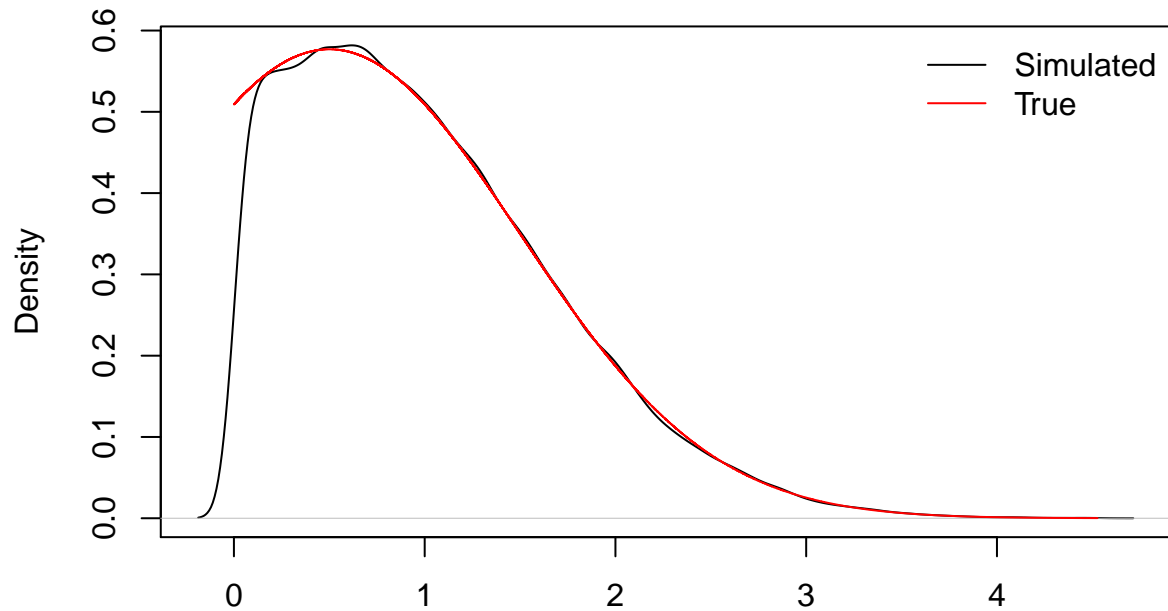
x<- rtn(n=100000,b=100,a=0,mu = 0.5,Sigma = 1)

mu = 0.5
Sigma = 1
b=100
a=0

x<-sort(x)
true_den <- dnorm((x-mu)/Sigma)/(pnorm((b-mu)/Sigma) - pnorm((a-mu)/Sigma))

plot(density(x),col = 1,main="X > 0")
lines(x,true_den,col = 2,type = "l")
legend('topright',legend=c('Simulated','True'),col=1:2,lty=1,bty='n')
```

**$X > 0$**



N = 100000 Bandwidth = 0.06248

Case  $x \in (-\infty, 0)$

```
x<- rtn(n=100000,b=0,a=-100,mu = 0.5,Sigma = 1)

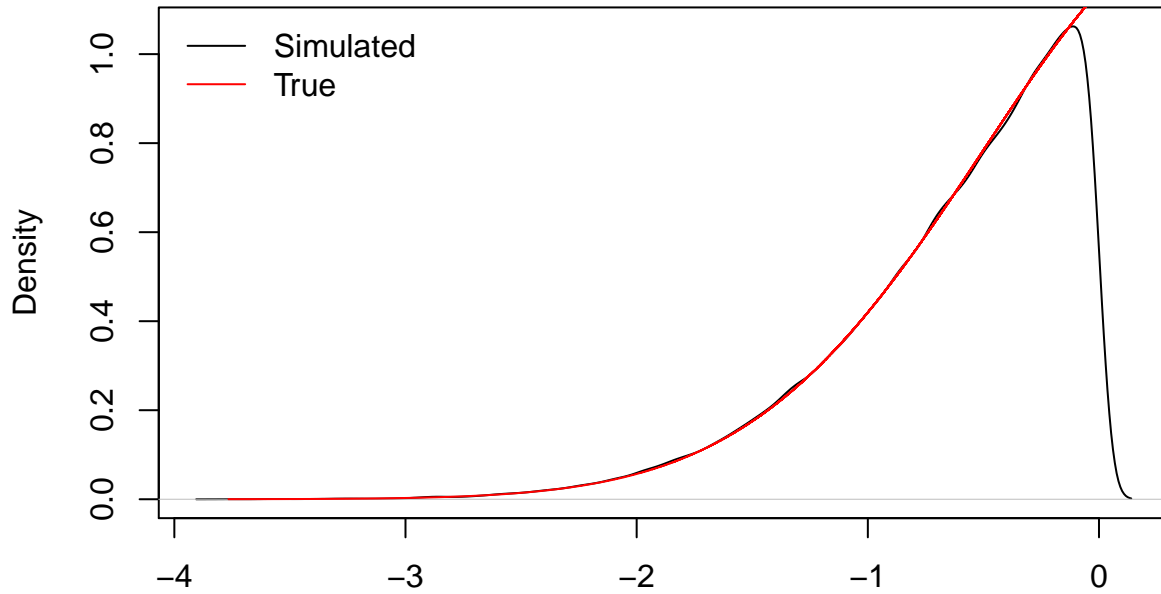
mu = 0.5
Sigma = 1
b=0
a=-100

x<-sort(x)
true_den <- dnorm((x-mu)/Sigma)/(pnorm((b-mu)/Sigma) - pnorm((a-mu)/Sigma))

plot(density(x),col = 1,main="X < 0")
lines(x,true_den,col = 2,type = "l")
legend('topleft',legend=c('Simulated','True'),col=1:2,lty=1,bty='n')
```



**$X < 0$**



**N = 100000 Bandwidth = 0.04657**

- c) Implement a Gibbs sampler to fit the same mixed model as fitted in Stan in a), but now with a probit link. As before, fit 4 chains, each running for 2000 iterations, with the first 1000 iterations discarded as burn-in. Perform graphical convergence checks and Gelman-Rubin diagnostics. Report posterior means, standard deviations and 95 % central credible intervals for  $\sigma, \beta, \mathbf{u}$  by combining chains.
- d) For the co-efficients  $\beta, \mathbf{u}$ , calculate the mean of the ratio of the posterior means  $\beta_{i,\text{logit}}/\beta_{i,\text{probit}}, \mathbf{u}_{i,\text{logit}}/\mathbf{u}_{i,\text{probit}}$  obtained when fitting the logistic mixed model and the probit mixed model. To do this, you will need to apply the `extract` function to the stan model object. Once calculated, multiply the iterations obtained assuming a probit link by this constant and compare to the iterations obtained assuming a logit link.
- e) The logistic link can be written in the same way as the probit link, but instead of  $e_i \sim \mathcal{N}(0, 1)$ , the error term is  $e_i \sim \text{Logistic}(0, 1)$ . By evaluating the standard normal and logistic inverse cdfs and superimposing the line  $y = mx$  where  $m$  is the posterior ratio, do you think the results in d) were surprising.